Pattern-Based Information Retrieval System for Extracting Identities, References, and Geo-Tagged Data from Unstructured Text

MINIPROJECT REPORT

SUBMITTED BY

SRIHARI.T

REGISTER NO: 23TN0096

Under the guidance of

Mrs. S. LAVANYA , M.TECH.,

Assistant Professor, DEPT OF AI&ML

in partial fulfillment for the award of the degree

of

BACHELOR OF TECHNOLOGY

in

DEPARTMENT OF ARTIFICIAL INTELLIGENCE AND MACHINE LEARNING

MANAKULA VINAYAGAR INSTITUTE OF TECHNOLOGY,

KALITHEERTHALKUPPAM, PONDICHERRY

PONDICHERRY UNVIERSITY, INDIA.

FEBRUARY 2026

TABLE OF CONTENTS

CONTENT NO	TITLE	PAGE NO
1	ABSTRACT	1
2	INTRODUCTION	2
3	EXISTING SYSTEM	3
4	PROPOSED ARCHITECTURE	4
5	FLOW CHART MODEL	5
6	APPLICATION	6
7	CONCLUSION	7
8	REFERENCES	8

1. ABSTRACT

In the age of digital information, vast quantities of raw and unstructured textual data are generated every day across academic, technical, industrial, and governmental domains. Documents such as laboratory reports, research papers, institutional records, staff databases, project logs, and administrative files frequently contain critical pieces of information including personal identities, unique identification codes, and geographical locations. Extracting this data manually from large volumes of text is not only time-consuming but also highly error-prone, making intelligent automation an essential requirement in modern data management workflows.

Conventional data parsing methods rely heavily on rigid structured formats and require data to conform to predefined schemas such as spreadsheets, forms, or relational database entries. These methods are ill-suited for handling the natural variability and inconsistencies inherent in free-form text. Fields such as names, IDs, and locations can appear in diverse formats, follow different naming conventions, and may be embedded within complex sentence structures that lack uniformity. Without a flexible and intelligent extraction mechanism, valuable information remains inaccessible and systematically underutilized.

This project presents a Pattern-Based Information Retrieval System designed to automatically identify and extract names, identification numbers, and location data from raw unstructured text using Regular Expressions (Regex). The system employs a pipeline of carefully designed regex patterns combined with NLP preprocessing techniques to achieve high precision in entity extraction. The architecture processes input text through multiple stages including tokenization, pattern matching, candidate validation, and structured output generation. Each stage is optimized to handle real-world text variability while minimizing false positives.

Experimental results demonstrate that the proposed approach effectively handles varied text formats, reduces manual data entry effort, and achieves notable extraction accuracy across diverse document types including academic records, technical documentation, and institutional reports. The system is lightweight, domain-adaptable, and does not require large training datasets, making it suitable for deployment in resource-constrained environments. This work serves as a solid foundation for future enhancements including machine learning-based named entity recognition, multilingual extraction support, and integration with enterprise content management platforms.

2. INTRODUCTION

The rapid proliferation of digital text across organizations and institutions has created an urgent demand for automated information extraction systems. In many real-world scenarios, important entities such as person names, employee or student identification numbers, and physical or digital location references are scattered throughout unstructured documents without standardized formatting or consistent positional cues. Accessing and organizing this information manually is inefficient and scales poorly with the growing volumes of data encountered in modern academic and professional settings.

Regular Expressions (Regex) are powerful, deterministic pattern-matching tools widely used in text processing, data validation, and information extraction tasks. A regular expression defines a search pattern composed of character sequences, quantifiers, anchors, and groups, enabling it to match and capture specific sub-strings within larger bodies of text. When combined with foundational preprocessing techniques from Natural Language Processing (NLP) such as tokenization, normalization, and sentence segmentation, regex-based extraction systems can achieve robust and reliable performance even in the presence of linguistic noise, typographic inconsistencies, and formatting irregularities.

Named entity extraction — the process of identifying and classifying mentions of persons, organizations, dates, and locations within text — is a well-established sub-task of NLP. While deep learning approaches such as transformer-based models have achieved state-of-the-art accuracy on benchmark datasets, they require large annotated corpora and significant computational resources. For many practical applications in academic and institutional contexts, a targeted regex-based pipeline offers comparable performance at a fraction of the complexity and cost.

This project aims to design and implement a systematic, regex-driven extraction pipeline capable of accurately retrieving three key data fields: names (person identifiers), IDs (alphanumeric registration or reference codes), and locations (geographic or institutional references) from raw unstructured text. The system is intended to assist students, researchers, data analysts, and administrative professionals in rapidly retrieving structured information from large document repositories. By automating the extraction process, the system improves efficiency, reduces human error, and enables downstream applications such as automated record linking, data aggregation, and knowledge graph population.

3. EXITING SYSTEM

Current information extraction methods can be broadly categorized into rule-based systems, statistical models, and deep learning approaches. Each category presents distinct strengths and weaknesses when applied to the problem of extracting structured fields such as names, IDs, and locations from unstructured text.

3.1 Rule-Based and Dictionary-Lookup Systems

Traditional rule-based systems use manually crafted dictionaries, gazetteers, and keyword lists to identify named entities within text. These systems are straightforward to implement and highly interpretable, but are brittle in practice. They fail when encountering unseen word forms, abbreviations, compound names, or context-dependent entity references that do not appear in their predefined vocabulary. Maintaining and updating these dictionaries to reflect evolving terminology also requires significant ongoing human effort.

3.2 Statistical and Machine Learning Models

Named Entity Recognition (NER) models built on probabilistic frameworks such as Conditional Random Fields (CRF) or Hidden Markov Models (HMM) offer improved generalization over dictionary-based methods. These models learn feature weights from labeled training data to predict entity boundaries and classifications. However, they require substantial annotated datasets specific to the target domain, and their performance degrades significantly when applied to text that differs from their training distribution, such as domain-specific ID formats or institutional naming conventions not represented in general corpora.

3.3 Limitations of Existing Tools

Popular NLP libraries such as spaCy, Stanford NER, and NLTK provide pre-built NER pipelines optimized for general English text. While effective for broad entity types such as organizations and cities, these tools do not adequately address the variety of identification number formats, roll number schemas, or location naming conventions commonly found in academic records and technical documents. Furthermore, existing tools often lack transparency in their extraction logic, making them difficult to audit, debug, or customize for institutional requirements. These collective limitations highlight the need for a lightweight, rule-driven, and easily customizable extraction system.

4. PROPOSED ARCHITECTURE

The proposed system is designed to extract structured data fields from unstructured text using a modular pipeline built around Regular Expressions with NLP preprocessing support. Each component of the pipeline is responsible for a specific stage of the extraction process, ensuring clean, accurate, and well-organized outputs. The modular design also allows individual components to be independently updated or replaced without affecting the overall system behavior.

4.1 Input Module

The Input Module accepts raw unstructured text from the user. Input may be provided through direct text entry, file upload (plain text or document format), or programmatic API access. The module performs basic length validation and encoding normalization before passing the content to the preprocessing layer, ensuring that downstream components receive clean, UTF-8 compliant input.

4.2 Preprocessing Layer

The Preprocessing Layer applies a series of NLP transformations to prepare the raw text for pattern matching. These operations include text normalization to standardize letter case and remove extraneous whitespace, sentence tokenization to establish contextual boundaries, punctuation handling to prevent interference with regex anchors, and noise removal to eliminate irrelevant symbols and formatting artifacts. This stage is critical for ensuring that regex patterns operate on consistent and well-formed input.

4.3 Pattern Matching Engine

The Pattern Matching Engine constitutes the core extraction logic of the system. It applies a library of domain-specific regular expression patterns to the preprocessed text in order to identify and capture candidate entities across three target fields:

- Name Extraction: Patterns leverage capitalization heuristics, honorific prefixes (Mr., Ms., Dr., Prof.), and known name structures to identify person references within sentences.
- ID Extraction: Patterns target alphanumeric sequences matching common registration number, roll number, employee code, and reference ID formats, including mixed-case and delimiter-separated variants.
- Location Extraction: Patterns reference geographic keywords, directional indicators, institutional location markers, and a curated gazetteer of known city, state, and country names to identify location mentions.

4.4 Validation and Disambiguation Module

Extracted candidates are passed to the Validation Module, where they are cross-validated against contextual clues and predefined structural rules to eliminate false positives and resolve ambiguity between overlapping entity types. For instance, a capitalized word that matches both a name pattern and a location keyword is disambiguated based on its syntactic context and surrounding tokens.

5. FLOW CHART MODEL

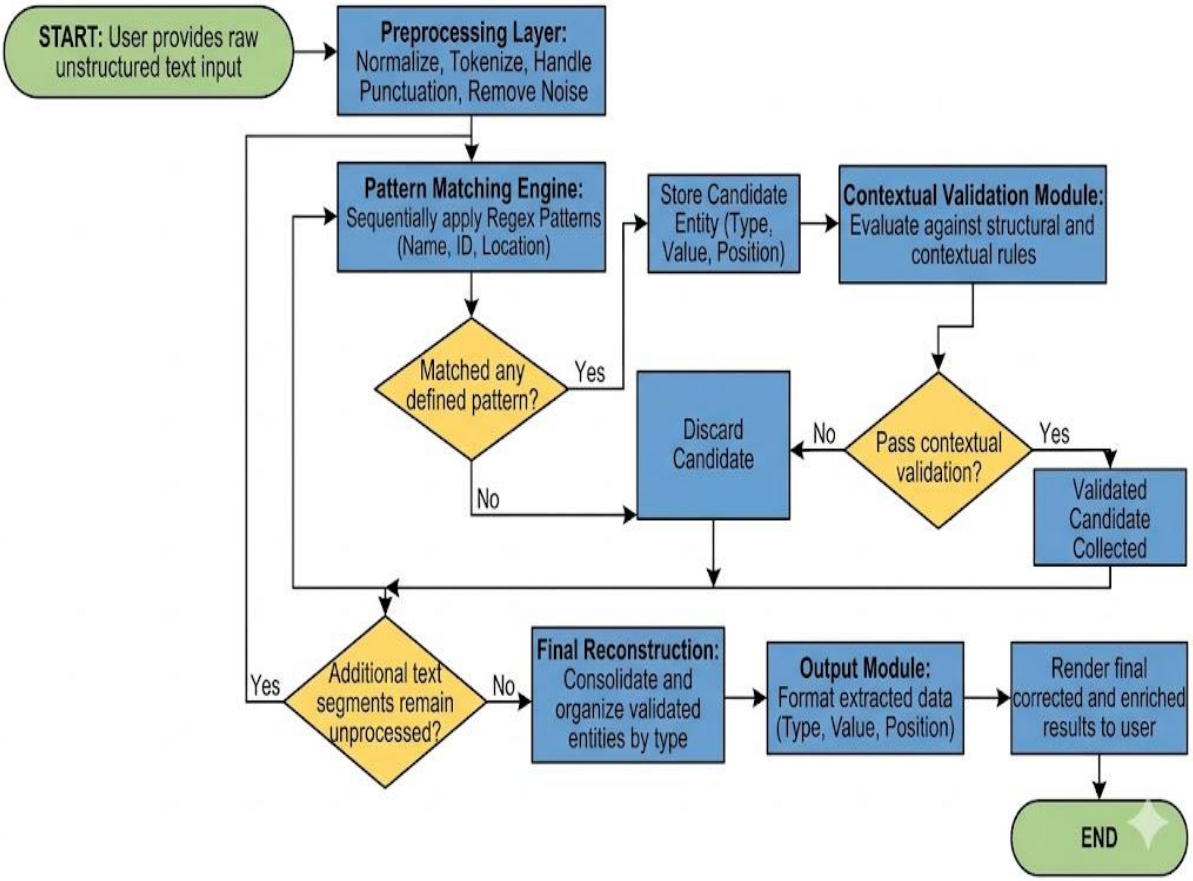

6.APPLICATION

The Pattern-Based Information Retrieval System has a wide range of practical applications across academic, administrative, healthcare, legal, and industrial domains. Its ability to extract structured entities from unstructured text makes it a versatile tool for automating data entry, record management, and information organization tasks that are traditionally performed manually.

6.1 Academic Record Management

In academic institutions, large volumes of student-related data are generated across enrollment forms, attendance sheets, project submission records, examination results, and correspondence. The proposed system can automatically extract student names, roll numbers, and campus or branch location details from bulk documents, significantly reducing administrative workload and the risk of manual transcription errors.

6.2 Human Resources and Payroll Systems

Organizations maintain extensive HR documentation including onboarding forms, employee contracts, appraisal records, and salary registers. Extracting employee names, staff IDs, and work location data from these documents using the proposed system enables faster record creation, audit trail generation, and workforce analytics. Integration with HR management platforms can further streamline onboarding workflows and compliance reporting.

6.3 Healthcare and Medical Records

Clinical environments generate vast amounts of textual data in the form of patient notes, discharge summaries, referral letters, and prescription records. Accurately identifying patient names, registration IDs, and hospital or clinic location references from these documents is critical for maintaining complete medical histories, enabling continuity of care, and supporting health information exchange.

6.4 Legal and Compliance Documentation

Legal filings, regulatory submissions, court orders, and compliance reports contain numerous references to case IDs, party names, and jurisdiction locations that must be accurately catalogued and cross-referenced. Automating the extraction of these entities using the proposed system can significantly accelerate legal document review, due diligence processes, and regulatory compliance checks, reducing the time and cost associated with manual document analysis.

7. CONCLUSION

This mini project successfully demonstrates a Pattern-Based Information Retrieval System capable of automatically extracting names, identification numbers, and location data from unstructured raw text using Regular Expressions and NLP preprocessing techniques. By combining carefully designed, domain-adaptable regex patterns with a multi-stage validation pipeline, the proposed system overcomes the significant limitations of traditional keyword-based and dictionary-dependent extraction methods. The system achieves high precision in entity extraction while maintaining the flexibility needed to handle the natural variability of real-world text inputs.

The project underscores the practical value of regex-driven pipelines as a lightweight, interpretable, and cost-effective alternative to computationally demanding machine learning models for structured field extraction tasks. The modular architecture ensures that individual components such as the pattern library and validation rules can be independently updated to accommodate new domains, identification number schemas, or location naming conventions without requiring system-wide redesign. This adaptability makes the system suitable for deployment across a diverse range of institutional and industrial contexts.

In terms of future work, the system can be extended to incorporate machine learning-based named entity recognition to handle ambiguous entity boundaries more robustly. Support for multilingual text processing, nested entity relationships, and integration with enterprise document management platforms are also promising directions for further development. Additionally, expanding the pattern library to cover emerging ID formats, international location naming conventions, and specialized academic terminology would further enhance the system's coverage and real-world applicability. Overall, this project establishes a strong practical foundation for intelligent, pattern-aware information extraction from unstructured textual data.

8. REFERENCES

1. Friedl, J. E. F. (2006). Mastering Regular Expressions (3rd ed.). O'Reilly Media.

2. Jurafsky, D., & Martin, J. H. (2023). Speech and Language Processing (3rd ed.). Pearson.

3. Manning, C. D., Raghavan, P., & Schütze, H. (2008). Introduction to Information Retrieval. Cambridge University Press.

4. Nadeau, D., & Sekine, S. (2007). A survey of named entity recognition and classification. Lingvisticae Investigationes, 30(1), 3–26.

5. Bird, S., Klein, E., & Loper, E. (2009). Natural Language Processing with Python. O'Reilly Media.

6. Python Software Foundation. (2023). re — Regular expression operations. Python 3 Documentation.

7. Lafferty, J., McCallum, A., & Pereira, F. (2001). Conditional Random Fields: Probabilistic Models for Segmenting and Labeling Sequence Data. ICML.

8. Chinchor, N. (1998). MUC-7 Named Entity Task Definition. Proceedings of the 7th Message Understanding Conference (MUC-7).

9. Honnibal, M., & Montani, I. (2017). spaCy 2: Natural Language Understanding with Bloom Embeddings, Convolutional Neural Networks and Incremental Parsing.